Journée

DEV:Ops

Des dashboards pour tous avec ELK

10 Juin 2014

Vincent Spiewak – @vspiewak

Xebia

Journée DEV:OPS

Introduction

Xebia

Speaker

Vincent Spiewak
@vspiewak

- 5 ans XP
- Master TA (UPMC)
- http://blog.xebia.fr
- @vspiewak

Agenda

- Introduction
- Logstash
- Monitoring Système
- Monitoring JMX
- Log As A Service
- Monitoring Métier / BI
- Cluster ELK
- Vagrant (démos)

Stack

Logstash

ETL

Elasticsearch

Stockage

Kibana

Visualisation

Logstash

Inputs	Codecs	Filters	Outputs
41	20	50	55
» stdin	» plain	» grok	» stdout
» file	» json	» date	» file
» udp	» line	» drop	» udp
» tcp	» multiline	» mutate	» tcp
» rabbitmq	» dots	» geoip	» rabbitmq
» s3	» msgpack	» anonymize	» elasticsearch
» …	» …	» …	» …

Elasticsearch

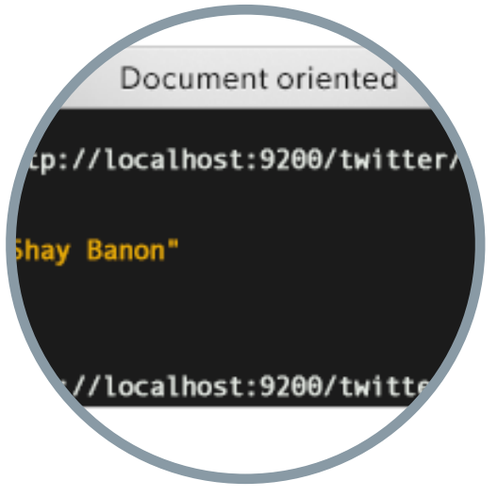

Document

Schema Free

Full Text

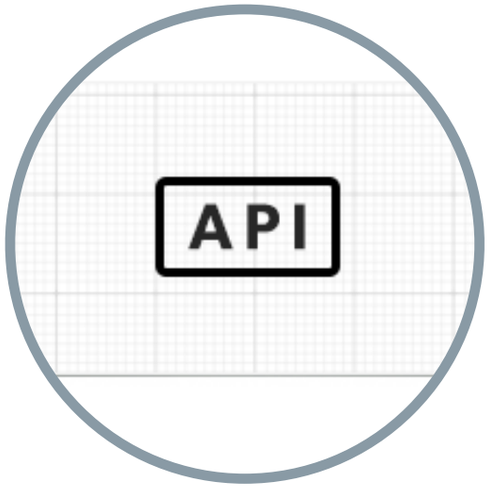

REST

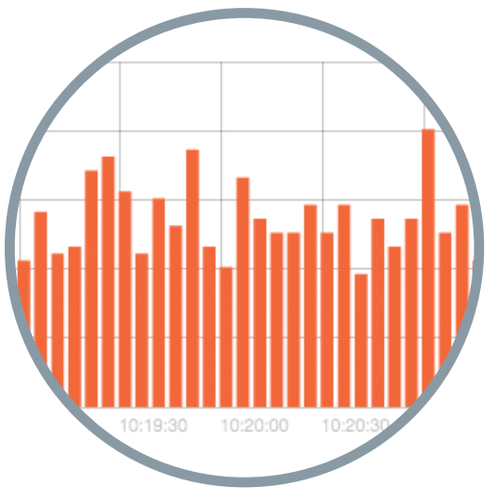

Real Time

Distributed

HA

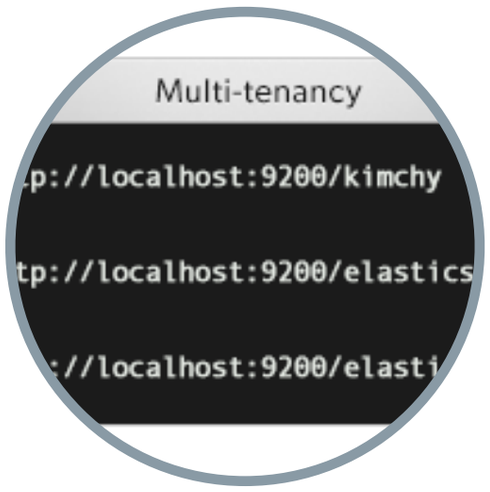

Multi-tenancy

Kibana


Logstash Search
Kibana 3 milestone pre-3
Query Alias
Images
Options
Search
Query
php
png OR gif
css
Chart
Zoom Out | Dynamic (12239) | Images (1
nt per 30m | (40354 hits)
800
600
400
200
0
16:00 07/24
20:00 07/24
00:00 07/25
04:00 07/25
08:00 07/25
12:00 07/25
16:00 07/25
20:00 07/25
00:00 07/26
04:00 07/26
08:00 07/26
12:00 07/26
Terms
Term Count Action
id 1856
br 1567
pk 1437
ng 1207
bd 1172
ru 1078
jp 959
mx 678
Missing field 9818
Other values 20582
id (1856) br (1567) pk (1437) ng (1207) bd (1172) ru (1078) jp (959) mx (678)
2000
1500
1000
500
0
filtering
time must
field : @timestamp
from : "2013-07-24T21:13:04.742Z"
to : "2013-07-26T21:13:04.743Z"
terms mustNot
field : country
value : cn
Toggle
terms mustNot
field : country
value : us
terms mustNot
field : country
value : id


HTML


B

Dashboards Adaptés

Infrastructure	Application	Business	Objectifs
» serveur	» stacktrace	» client	» résultat
» charge	» warn, error	» produit	» progression

Architecture

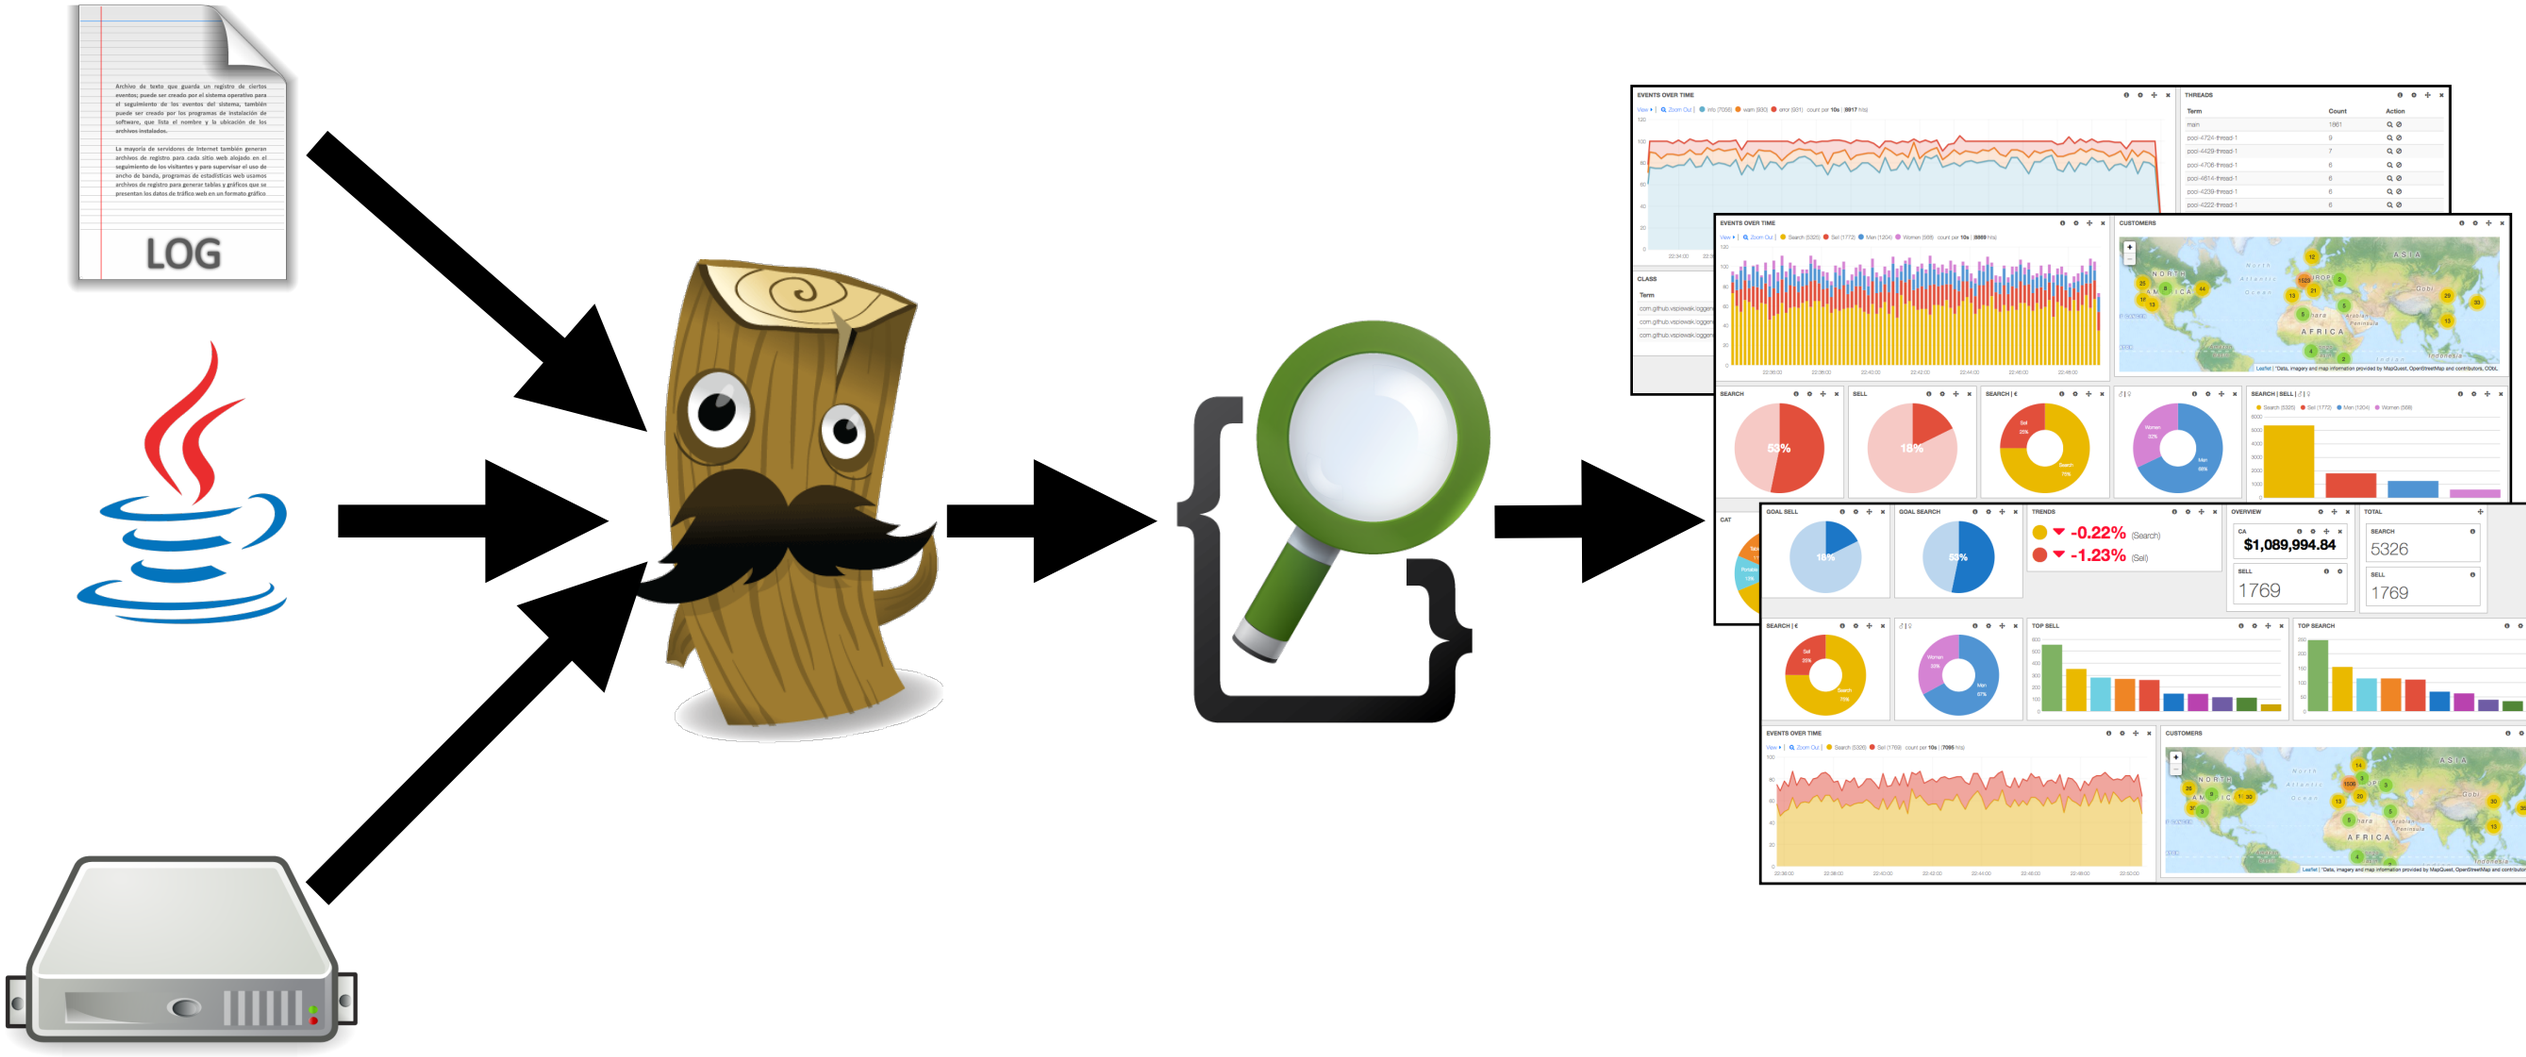

Journée DEV:OPS

Logstash

Xebia

Logstash – Entrée/Sortie

```
input {
  stdin {}
}

# filters


output {
  stdout { codec => json }
}
```

Logstash – Run

```
$ java -jar logstash.jar agent –f app.conf

2011-04-19T03:44:01.103Z 55.3.244.1 GET /index.html 15824 0.043

{
        "message" => "2011-04-19T03:44:01.103Z GET /index.html 15824 0.043",
     "@timestamp" => "2013-11-03T19:48:53.175Z",
       "@version" => "1",
           "host" => "macbook"
}
```

Logstash – Patterns

https://github.com/logstash/logstash/blob/master/patterns

```
USERNAME [a-zA-Z0-9._-]+

USER %{USERNAME}

INT (?:[+-]?(?:[0-9]+))

WORD \b\w+\b

NOTSPACE \S+

DATA .*?

GREEDYDATA .*

HTTPDATE %{MONTHDAY}/%{MONTH}/%{YEAR}:%{TIME} %{INT}

COMBINEDAPACHELOG %{IPORHOST:clientip} …
```

Logstash – Filtre Grok

```
2011-04-19T03:44:01.103Z 55.3.244.1 GET /index.html 15824 0.043
```

```
filter {
   grok {
      match =>

            [ "message",
              "%{TIMESTAMP_ISO8601:date} %{IP:client} %{WORD:method}
              %{URIPATHPARAM:uri} %{NUMBER:bytes} %{NUMBER:duration}"
            ]
   }
}
```

Filtre Grok – Sortie

```
2011-04-19T03:44:01.103Z 55.3.244.1 GET /index.html 15824 0.043
```

```
{
    "@timestamp" => "2013-12-01T21:19:11.303Z",
     "@version" => "1",
        "@bytes" => "15824",
      "@client" => "55.3.244.1",
          "date" => "2011-04-19T03:44:01.103Z",
    "@duration" => "0.043",
          "host" => "macbookpro",
       "message" => "2011-04-19T03:44:01.103Z 55.3.244.1 GET /index.html 15824 0.043",
        "method" => "GET",
           "uri" => "/index.html",
}
```

Filtre Date – @Timestamp

```
filter {
    date {
        match => [ "date", "ISO8601" ],
    }
}
```

Filtre Date – @Timestamp

```
{
    "@timestamp" => "2011-04-19T03:44:01.103Z",
      "@version" => "1",
        "@bytes" => "15824",
       "@client" => "55.3.244.1",
          "date" => "2011-04-19T03:44:01.103Z",
     "@duration" => "0.043",
          "host" => "macbookpro",
       "message" => "2011-04-19T03:44:01.103Z 55.3.244.1 GET /index.html 15824 0.043",
        "method" => "GET",
           "uri" => "/index.html",
}
```

Sortie Elasticsearch


```
{
    _index: "logstash-2011.04.19",
    _type: "logs",
    _id: "VhSmjXzTQkyvFCHtwHVQVg",
    _version: 1,
    _score: null,
    _source: {
        message: "2011-04-19T03:44:01.103Z 55.3.244.1 GET /index.html 15824 0.043",
        @timestamp: "2011-04-19T03:44:01.103Z",
        @version: "1",
        host: "macbookpro",
        date: "2011-04-19T03:44:01.103Z",
        client: "55.3.244.1",
        method: "GET",
        request: "/index.html",
        bytes: "15824",
        duration: "0.043"
    },
    sort: [
        1303184641103
    ]
}
```

Filtres

- ajout d'un champ / type / tag
- suppression d'un champ
- split d'un champ
- conversion de type (string, int, float)
- IP => géolocation
- UA => device, browser, os, versions
- conditions
- etc…

Logstash – Sortie Elasticsearch

- host
- port
- cluster
- index => "logstash-%{+YYYY.MM.dd}"
- protocol
- ...

Journée DEV:OPS

Monitoring Système

Xebia

Monitoring Système: Collectd

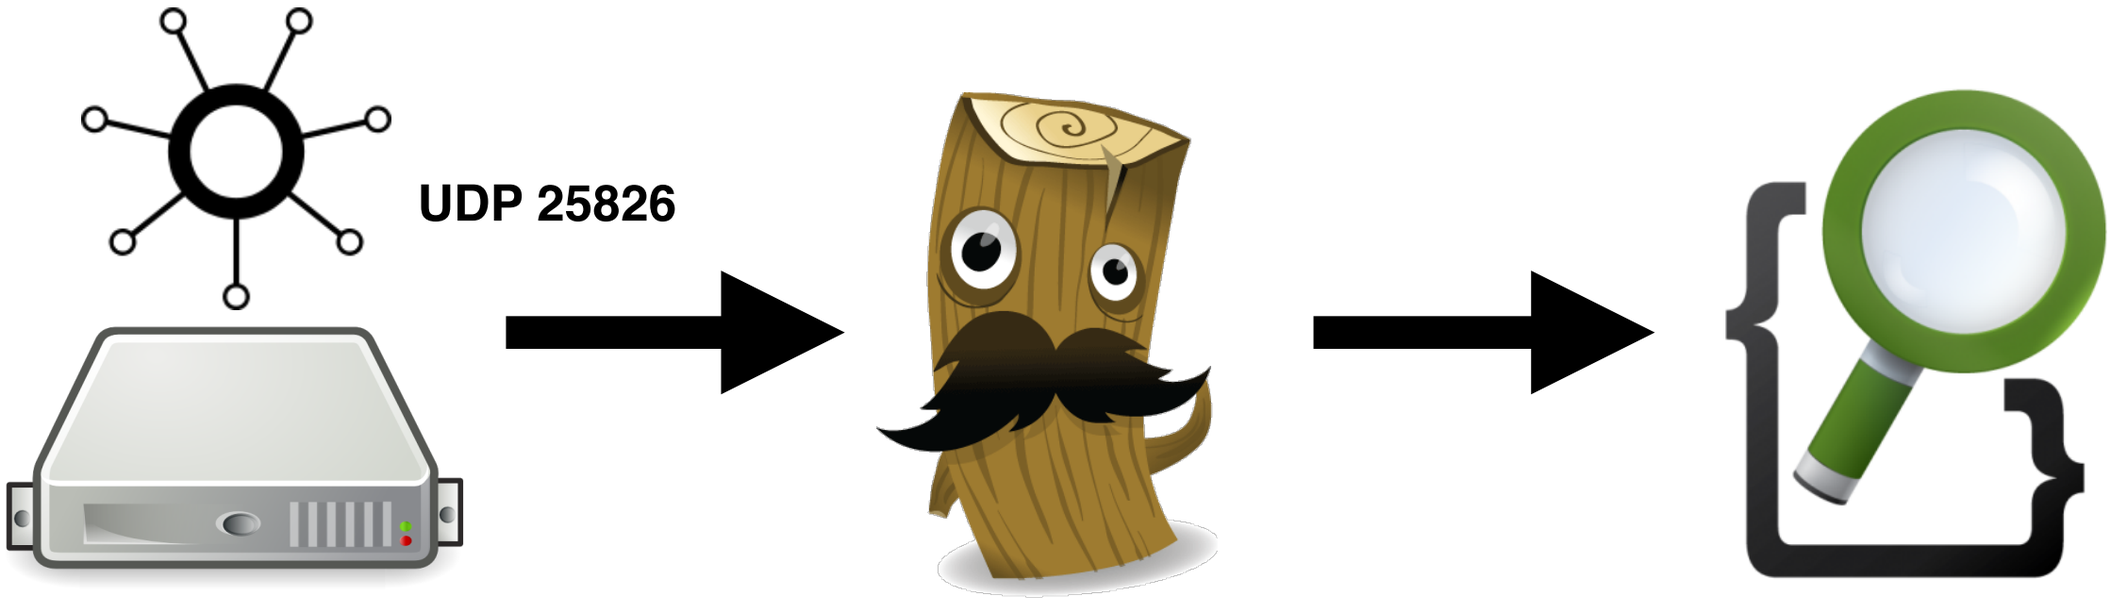

Collectd – Plugins

- cpu (jiffies)
- process
- users
- disk
- memory
- swap
- network
- Java / JMX
- MySQL
- …

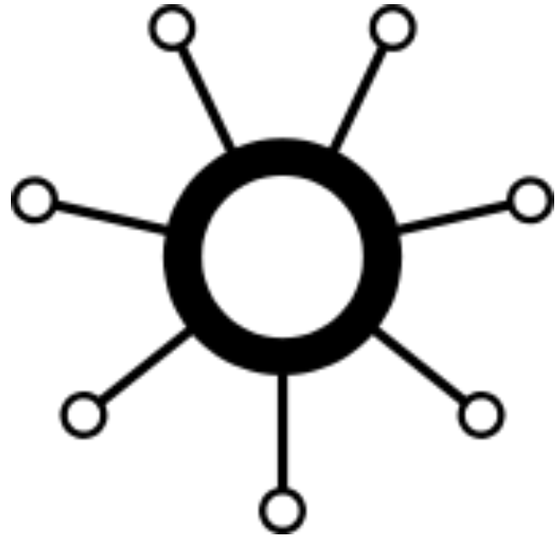

Collectd – Configuration

```
FQDNLookup true
LoadPlugin syslog
<Plugin syslog>
 LogLevel info
</Plugin>
LoadPlugin cpu
LoadPlugin df
LoadPlugin disk
LoadPlugin entropy
LoadPlugin interface
LoadPlugin irq
LoadPlugin load
LoadPlugin memory
LoadPlugin network
LoadPlugin processes
LoadPlugin rrdtool
LoadPlugin swap
LoadPlugin users
<Plugin interface>
Interface "eth0"
IgnoreSelected false
</Plugin>
<Plugin network>
    <Server "127.0.0.1" "25826">
    </Server>
</Plugin>
<Plugin rrdtool>
DataDir "/var/lib/collectd/rrd"
</Plugin>
Include "/etc/collectd/filters.conf"
Include "/etc/collectd/thresholds.conf"
```

Collectd – Logstash Conf

```
input {
  collectd {
     host => "127.0.0.1"
  }
}

output {
     elasticsearch {}
}
```

Elasticsearch – Samples

```
{
    "@version": "1",
    "@timestamp": "2014-06-09T23:01:11.000Z",
    "host": "precise64",
    "plugin": "memory",
    "collectd_type": "memory",
    "type_instance": "cached",
    "value": 267845632
}
```

```
{
    "@version": "1",
    "@timestamp": "2014-06-09T23:01:11.000Z",
    "host": "precise64",
    "plugin": "memory",
    "collectd_type": "memory",
    "type_instance": "used",
    "value": 703348736
}
```

Monitoring Système

system-survey

Démo

Monitoring Système

system-survey

Journée DEV:OPS

Monitoring JMX

Xebia

Monitoring JMX: Collectd JMX

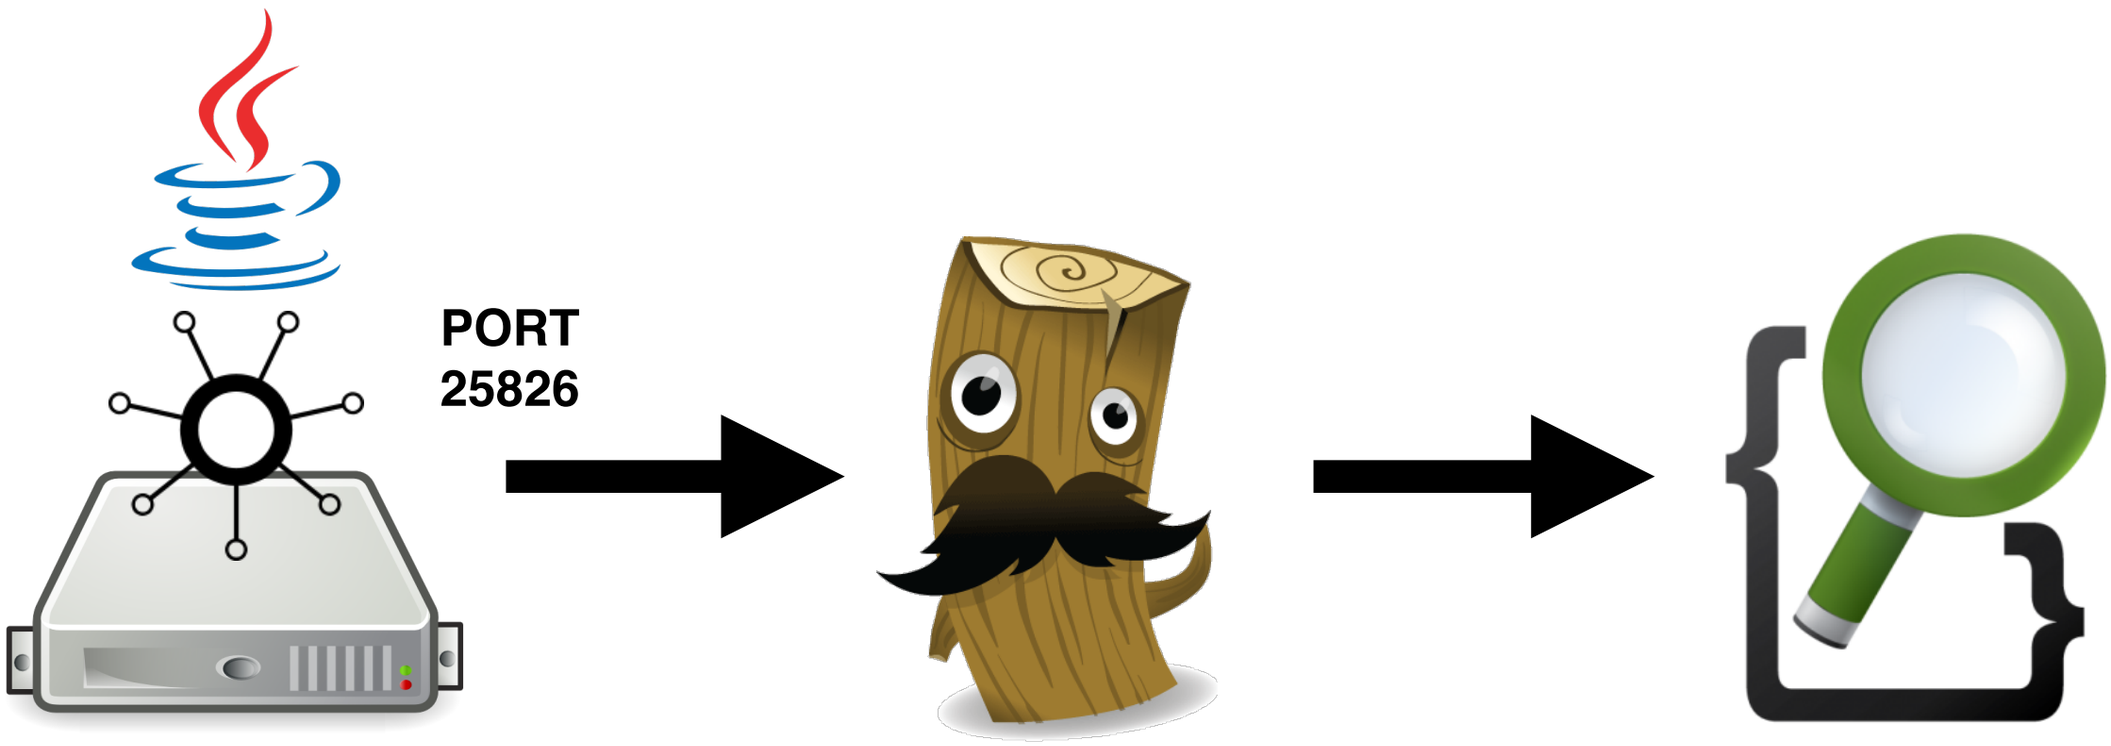

Collectd – Setup Java & JMX

```
# check dynamic libraries

ldd /usr/lib/collectd/java.so
```

```
# fix libjvm.so not found error
ln -s /usr/lib/jvm/java-7-openjdk-amd64/jre/lib/amd64/server/libjvm.so /usr/lib/libjvm.so
```

JConsole – SystemCpuLoad

pid: 14008 log-generator.jar -n 10 -r 1000 -e

Overview | Memory | Threads | Classes | VM Summary | MBeans

- JMImplementation
- com.sun.management
- java.lang
  - ClassLoading
  - Compilation
  - GarbageCollector
  - Memory
  - MemoryManager
  - MemoryPool
  - OperatingSystem
    - Attributes
      - OpenFileDescriptorCount
      - MaxFileDescriptorCount
      - CommittedVirtualMemorySize
      - TotalSwapSpaceSize
      - FreeSwapSpaceSize
      - ProcessCpuTime
      - FreePhysicalMemorySize
      - TotalPhysicalMemorySize
      - SystemCpuLoad
      - ProcessCpuLoad
      - Arch
      - Version
      - SystemLoadAverage
      - AvailableProcessors
      - Name
      - ObjectName
  - Runtime
  - Threading
- java.nio
- java.util.logging

Attribute value

Name	Value
SystemCpuLoad	0.8076923076923077

Refresh

MBeanAttributeInfo

Name	Value
Attribute:	
Name	SystemCpuLoad
Description	SystemCpuLoad
Readable	true
Writable	false
Is	false
Type	double

Descriptor

Name	Value
Attribute:	
openType	javax.management.openmbean.SimpleType(name=java.lang.Double)
originalType	double

Collectd – Plugin Java & JMX

```
<Plugin "java">
 JVMARG "-Djava.class.path=/usr/share/collectd/java/collectd-api.jar:/usr/share/collectd/java/generic-jmx.jar"
  LoadPlugin "org.collectd.java.GenericJMX"
  <Plugin "GenericJMX">

   <MBean "os">
      ObjectName "java.lang:type=OperatingSystem"

      <Value>
        Type "gauge"
        InstancePrefix "system_cpu_load"
        Attribute "SystemCpuLoad"
      </Value>

    </MBean>

    <Connection>
      ServiceURL "service:jmx:rmi:///jndi/rmi://localhost:9010/jmxrmi"
      Collect "os"
   </Connection>

  </Plugin>
</Plugin>
```

Elasticsearch – Samples


```
{
    "@version": "1",
    "@timestamp": "2014-06-09T23:01:11.000Z",
    "host": "localhost",
    "plugin": "GenericJMX",
    "collectd_type": "gauge",
    "type_instance": "system_cpu_load",
    "value": 0.558783783783838
}
```

JConsole – HeapMemoryUsage

pid: 14008 log-generator.jar -n 10 -r 1000 -e

Overview | Memory | Threads | Classes | VM Summary | MBeans

- JMImplementation
- com.sun.management
- java.lang
  - ClassLoading
  - Compilation
  - GarbageCollector
  - Memory
    - Attributes
      - ObjectPendingFinalizationCo
      - HeapMemoryUsage
      - NonHeapMemoryUsage
      - Verbose
      - ObjectName
    - Operations
    - Notifications
  - MemoryManager
  - MemoryPool
  - OperatingSystem
  - Runtime
  - Threading
- java.nio
- java.util.logging

Attribute value

Name	Value
HeapMemoryUsage	< Tabular Navigation > << < Composite Navigation >

Name	Value
commit...	15794176
init	16309632
max	504823808
used	10236760

Refresh

MBeanAttributeInfo

Name	Value
Attribute:	
Name	HeapMemoryUsage
Description	HeapMemoryUsage
Readable	true
Writable	false
Is	false
Type	javax.management.openmbean.CompositeData

Collectd – JMX – Type Table

```
# Heap memory usage
<MBean "memory-heap">
  ObjectName "java.lang:type=Memory"
  #InstanceFrom ""
  InstancePrefix "memory-heap"

  # Creates four values: committed, init, max, used
  <Value>
    Type "jmx_memory"
    Table true
    Attribute "HeapMemoryUsage"
  </Value>
</MBean>
```

Collectd – Types

/usr/share/collectd/types.db

```
gauge                   value:GAUGE:U:U
load                    shortterm:GAUGE:0:100, midterm:GAUGE:0:100, longterm:GAUGE:0:100
percent                 percent:GAUGE:0:100.1
jmx_memory              value:GAUGE:0:U
```

Elasticsearch – Samples


```
{
    "@version": "1",
    "@timestamp": "2014-06-09T23:01:11.000Z",
    "host": "localhost",
    "plugin": "GenericJMX",
    "plugin_instance": "memory-heap",
    "collectd_type": "jmx_memory",
    "type_instance": "used",
    "value": 62282808
}

{
    "@version": "1",
    "@timestamp": "2014-06-09T23:01:11.000Z",
    "host": "localhost",
    "plugin": "GenericJMX",
    "plugin_instance": "memory-heap",
    "collectd_type": "jmx_memory",
    "type_instance": "init",
    "value": 104857600
}
```

Collectd – Custom MBean

```
<MBean "flume-source">
  ObjectName "org.apache.flume.source:type=source-1"
  InstancePrefix "flume-source-1"
 <Value>
    Type "gauge"
    InstancePrefix "event_received_count"
    Attribute "EventReceivedCount"
  </Value>
  <Value>
    Type "gauge"
    InstancePrefix "event_accepted_count"
    Table false
    Attribute "EventAcceptedCount"
  </Value>
</MBean>
```

Elasticsearch – Samples

```
{
    "@version": "1",
    "@timestamp": "2014-06-09T23:09:41.000Z",
    "host": "localhost",
    "plugin": "GenericJMX",
    "plugin_instance": "flume-source-1",
    "collectd_type": "gauge",
    "type_instance": "event_accepted_count",
    "value": 1246501
}
```

Monitoring JVM / JMX

Flume JMX

Démo

Monitoring JVM / JMX

Flume JMX

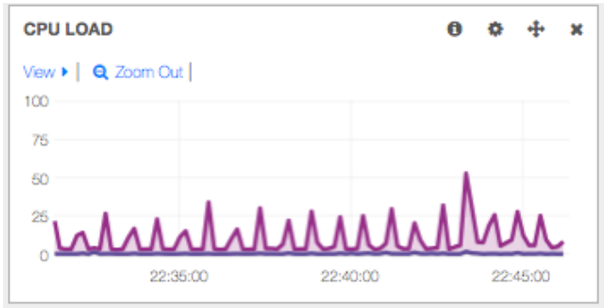

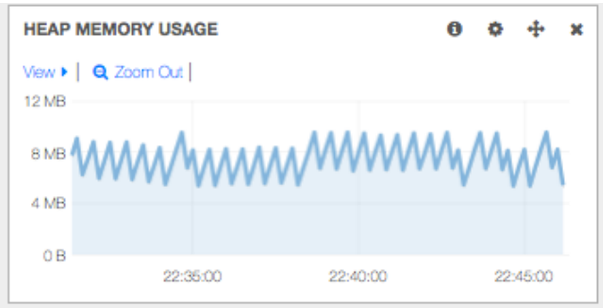

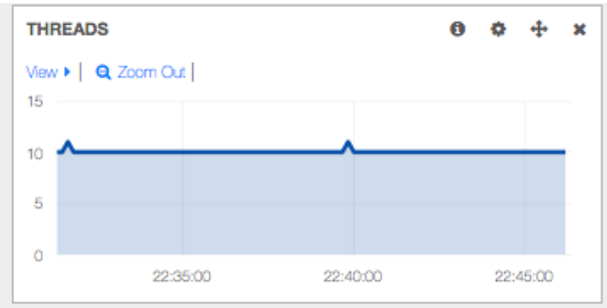

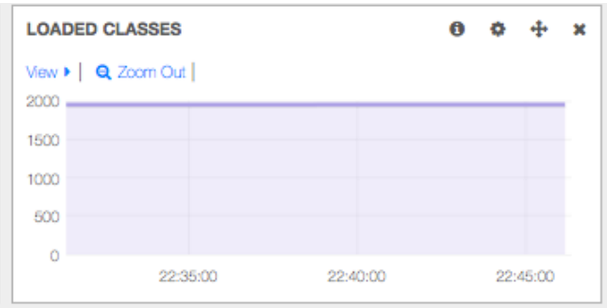

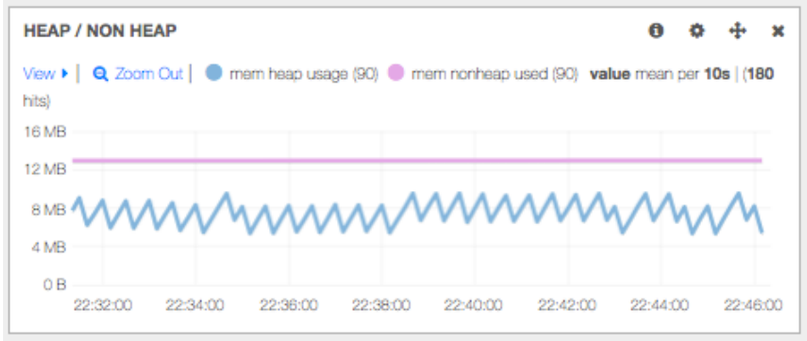

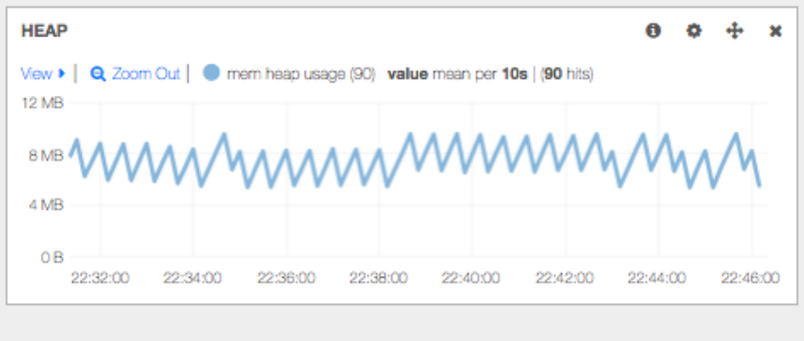

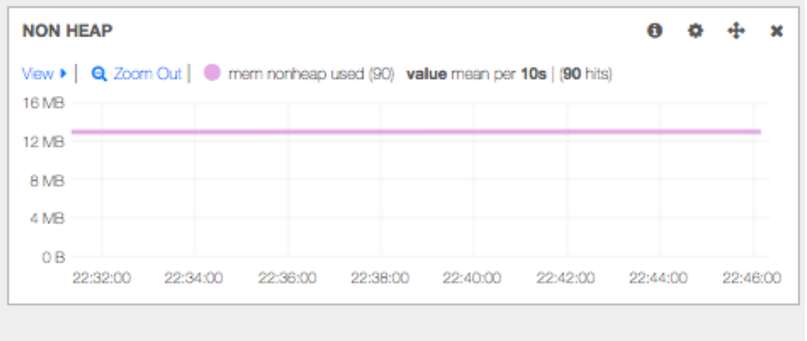

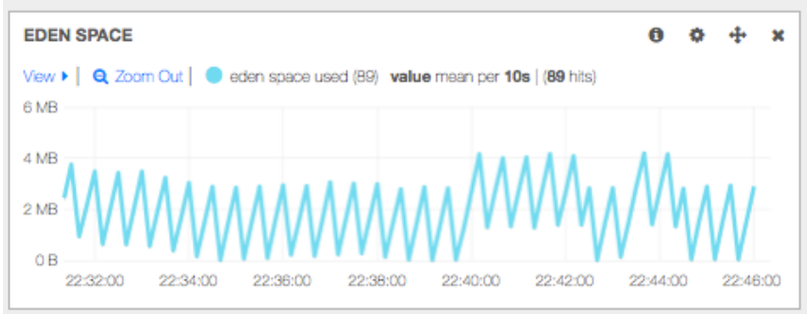

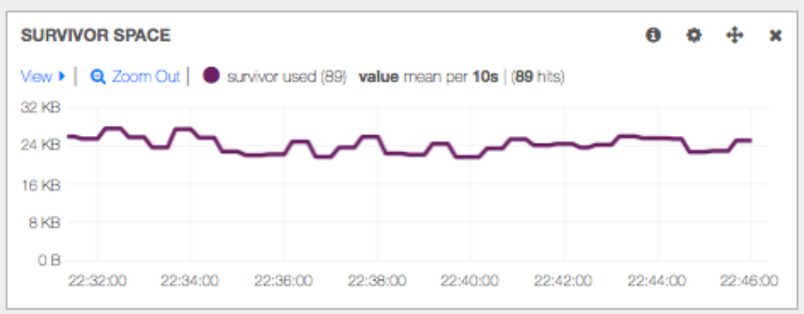

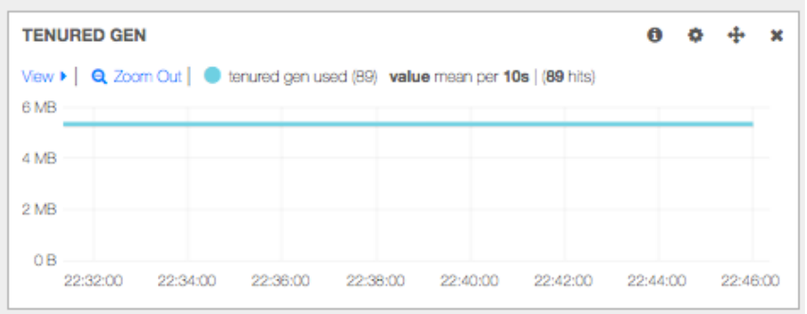

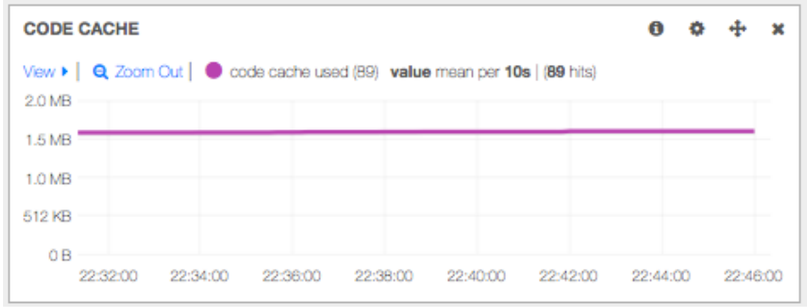

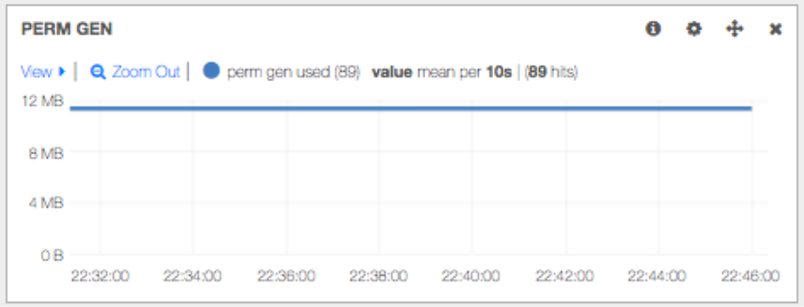

Journée DEV:OPS

Log As Service

Xebia

Log As A Service: SyslogAppender

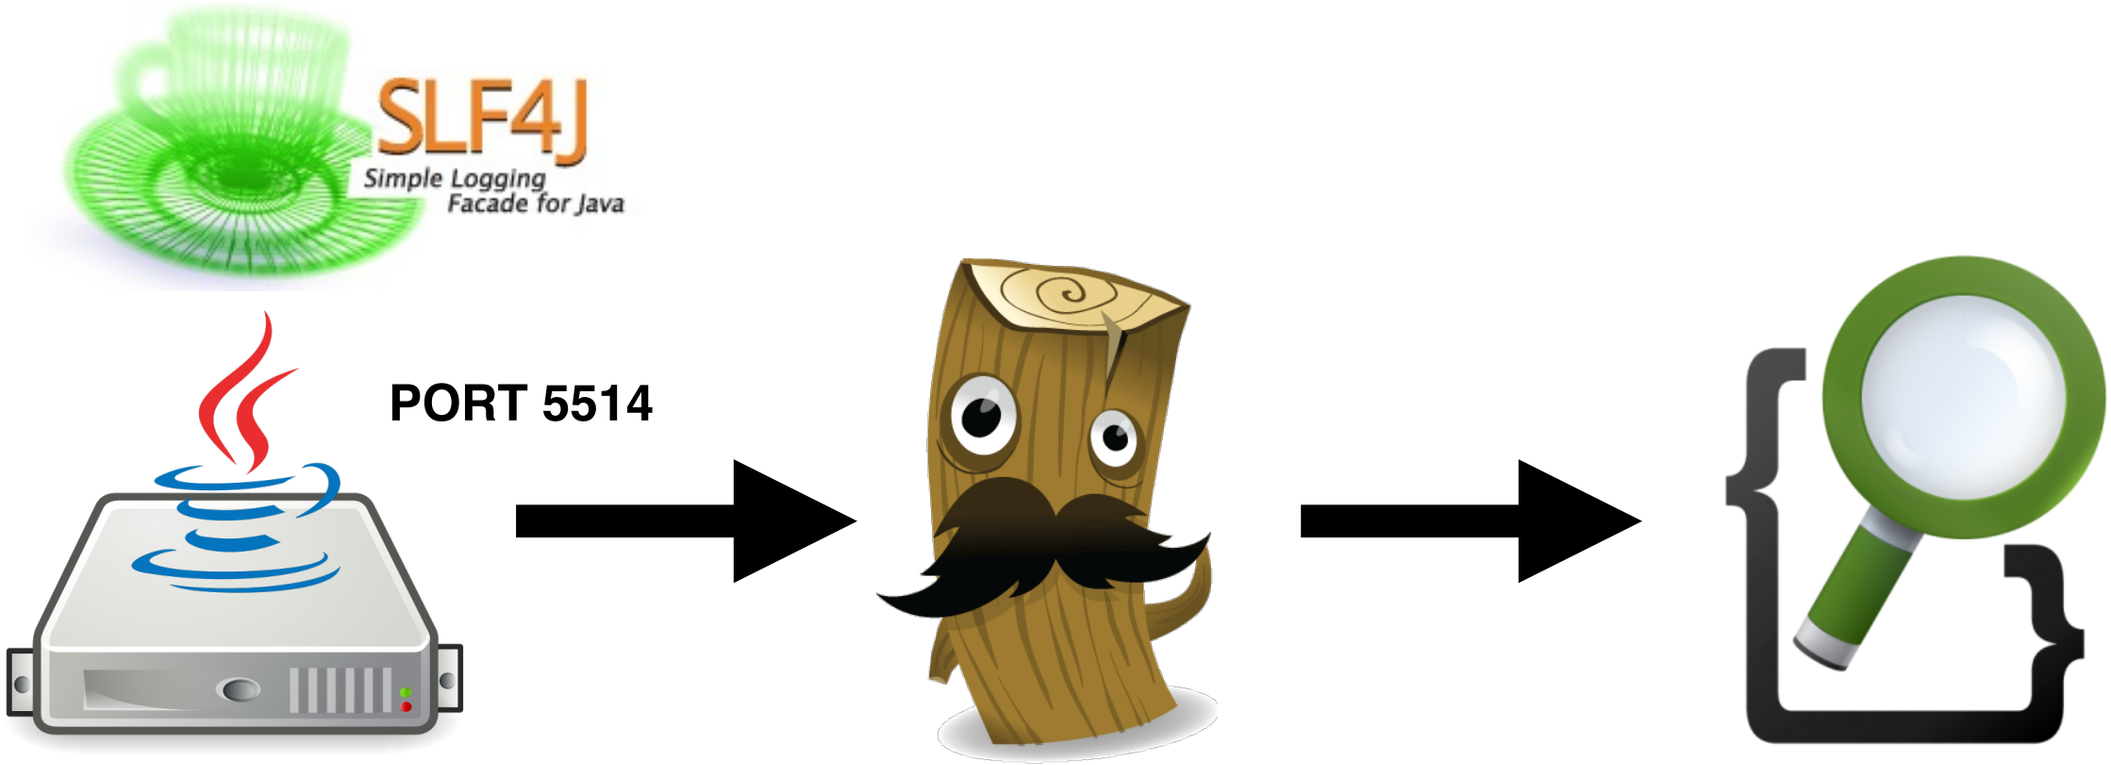

Log As A Service – Logback

SyslogAppender

```
<appender name="syslog" class="ch.qos.logback.classic.net.SyslogAppender">
  <syslogHost>127.0.0.1</syslogHost>
  <port>5514</port>
  <facility>user</facility>
  <suffixPattern>%d{dd-MM-yyyy HH:mm:ss.SSS} [%thread] %level %logger - %msg%n</suffixPattern>
</appender>
```

Logstash – Syslog configuration

```
input {
  udp {
     port => "5514"
  }
}

filter {
   grok {
      patterns_dir => "./patterns"
      match => ["message","%{LOGBACK_SYSLOG}"]
   }
}
```

```
filter {
   date {
      match => ["log_date","dd-MM-YYYY HH:mm:ss.SSS"]
   }
}

output {
  elasticsearch {}
}
```

Logstash – Pattern

```
LOG_DATE %{MONTHDAY}-%{MONTHNUM}-%{YEAR} %{HOUR}:%{MINUTE}:%{SECOND}.[0-9]{3}

SYSLOG_BASE %{SYSLOG5424PRI}%{SYSLOGTIMESTAMP:syslog_timestamp} %{SYSLOGHOST:syslog_host}
SYSLOG %{SYSLOG_BASE} %{GREEDYDATA:syslog_message}

LOGBACK_SYSLOG_BASE %{SYSLOG_BASE} %{LOG_DATE:log_date} \[%{NOTSPACE:thread}\] %{LOGLEVEL:log_level} %{NOTSPACE:classname}
LOGBACK_SYSLOG %{LOGBACK_SYSLOG_BASE} %{GREEDYDATA:log_msg}
```

Monitoring Log
Syslog

Démo

Monitoring Log
Syslog

EVENTS OVER TIME

View ▸ | Zoom Out | info (7056) warn (930) error (931) count per **10s** | (**8917** hits)

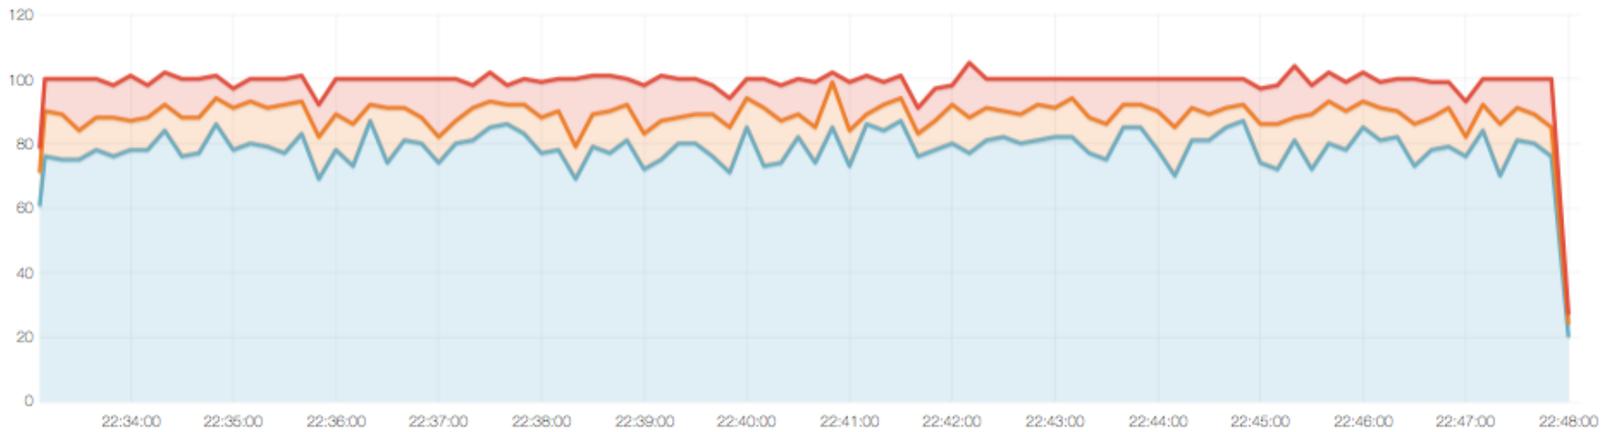

THREADS

Term	Count	Action
main	1861	
pool-4724-thread-1	9	
pool-4429-thread-1	7	
pool-4706-thread-1	6	
pool-4614-thread-1	6	
pool-4239-thread-1	6	
pool-4222-thread-1	6	
pool-4154-thread-1	6	
pool-4834-thread-1	5	
pool-4820-thread-1	5	

CLASS

Term	Count	Action
com.github.vspiewak.loggenerator.SearchRequest	5274	
com.github.vspiewak.loggenerator.App	1861	
com.github.vspiewak.loggenerator.SelRequest	1782	

LEVEL

Term	Count	Action
INFO	7056	
ERROR	931	
WARN	930	

CLASS

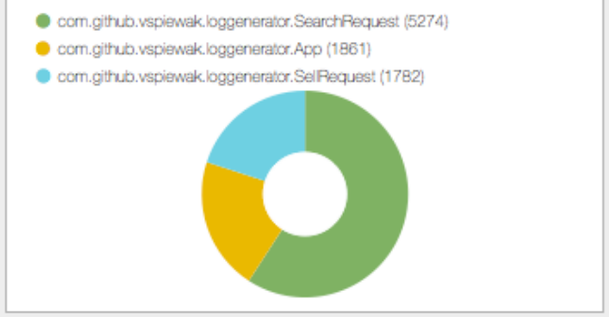

LEVEL

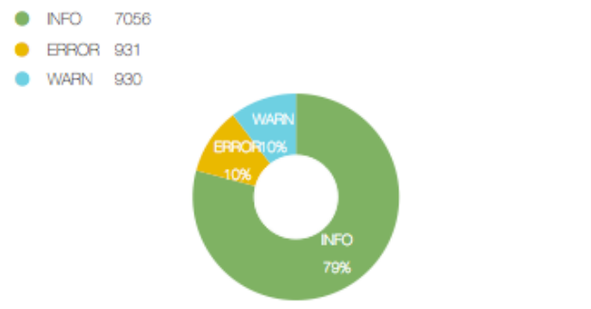

Journée DEV:OPS

Monitoring Métier– BI

Xebia

GeekShop
Problème

- Quels sont les produits les plus achetés ?
- Quelle est la répartition H/F de mes clients ?
- Quels sont mes clients les plus fidèles ?
- Combien de femmes à Paris ont acheté un iPod Touch Bleu 32 Go entre le 12 octobre 2012 à 14h30 et le 4 novembre 2013 à 19h ?

GeekShop – Format Logs

```
09-06-2014 21:27:42.228 [pool-32-thread-1] INFO
com.github.vspiewak.loggenerator.SearchRequest -
id=317&ua=Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/534.24
(KHTML, like Gecko) Chrome/11.0.696.65 Safari/
534.24&ip=94.228.34.210&category=Mobile
```

```
09-06-2014 21:27:42.227 [pool-32-thread-1] INFO
com.github.vspiewak.loggenerator.SellRequest - id=313&ua=Mozilla/
5.0 (Windows; U; Windows NT 6.1; en-US; rv:1.9.2.17) Gecko/
20110420 Firefox/
3.6.17&ip=202.46.52.35&email=client314@gmail.com&sex=M&brand=Appl
e&name=iPod Touch&model=iPod Touch - Jaune - Disque
32Go&category=Baladeur&color=Jaune&options=Disque
32Go&price=329.0
```

GeekShop – Après Logstash


```
{
    "_index": "logstash-2014.06.09",
    "_type": "app-log",
    "_id": "gaQXRn9mROiAGjhBZ2h2Og",
    "_version": 1,
    "found": true,
    "_source": {
        "message": "09-06-2014 21:27:42.228 [pool-32-thread-1] INFO com.github.vspiewak.loggenerator.SearchRequest - id=317&ua=Mozilla/5.0 (X11;
Linux x86_64) AppleWebKit/534.24 (KHTML, like Gecko) Chrome/11.0.696.65 Safari/534.24&ip=94.228.34.210&category=Mobile",
        "@version": "1",
        "@timestamp": "2014-06-09T19:27:42.228Z",
        "type": "app-log",
        "host": "precise64",
        "path": "/home/vagrant/app.log",
        "log_date": "09-06-2014 21:27:42.228",
        "thread": "pool-32-thread-1",
        "log_level": "INFO",
        "classname": "com.github.vspiewak.loggenerator.SearchRequest",
        "log_msg": "- id=317&ua=Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/534.24 (KHTML, like Gecko) Chrome/11.0.696.65 Safari/
534.24&ip=94.228.34.210&category=Mobile",
        "id": 317,
        "ua": "Mozilla/5.0 (X11; Linux x86_64) AppleWebKit/534.24 (KHTML, like Gecko) Chrome/11.0.696.65 Safari/534.24",
        "ip": "94.228.34.210",
        "category": "Mobile",
        "tags": [
            "search"
        ],
        "geoip": {
            "ip": "94.228.34.210",
            "country_code2": "GB",
            "country_code3": "GBR",
            "country_name": "United Kingdom",
            "continent_code": "EU",
            "latitude": 51.5,
            "longitude": -0.12999999999999545,
            "timezone": "Europe/London",
            "location": [
                -0.12999999999999545,
                51.5
            ]
        },
        "useragent": {
            "name": "Chrome",
            "os": "Linux",
            "os_name": "Linux",
            "device": "Other",
            "major": "11",
            "minor": "0",
            "patch": "696"
        }
    }
}
```

Kibana: Terms & Analysers

TOP CUSTOMERS

Term	Count	Action
gmail.com	367	
client920	2	
client828	2	
client706	2	
client677	2	

TOP PRODUCTS

Term	Count	Action
ipod	153	
touch	78	
iphone	72	
macbook	31	
nano	19	

Elasticsearch Template Mapping

Change analyser on specific indexes & fields

```
curl -XPUT http://localhost:9200/_template/logstash_per_index -d '{
    "template" : "logstash*",
    "mappings" : {
        "_default_" : {
           "properties" : {
                "@timestamp": { "type": "date", "index": "not_analyzed" },
                "ip": { "type" : "ip", "index": "not_analyzed" },
                "name": { "type" : "string", "index": "not_analyzed" },
                "options": { "type" : "string", "index": "not_analyzed" },
                "email": { "type" : "string", "index": "not_analyzed" }
           }
        }
    }
}'
```

Monitoring Métier / Business

eshop-survey

Démo

Monitoring Métier / Business

eshop-survey

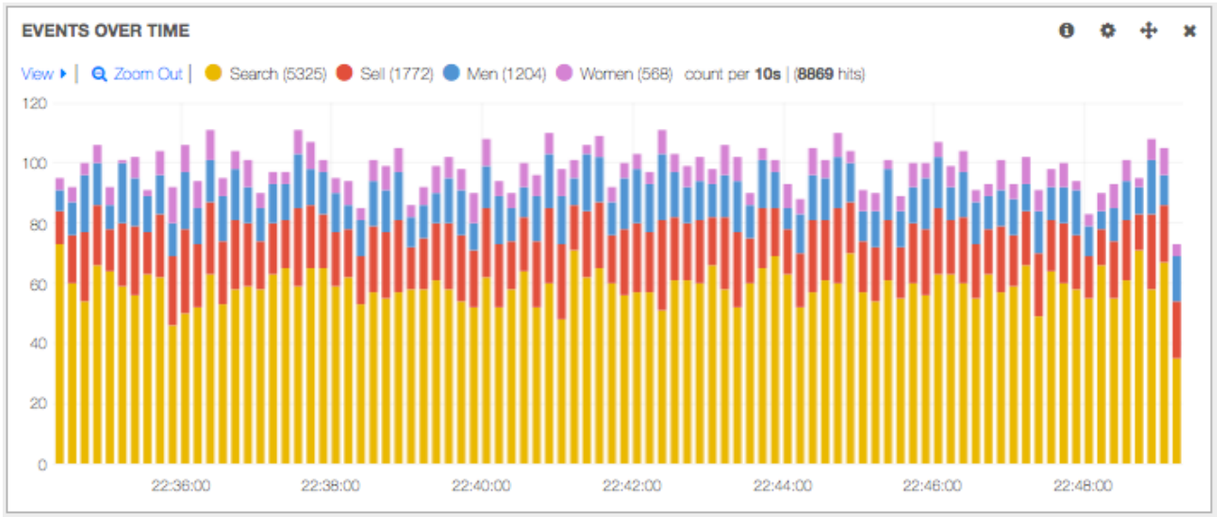

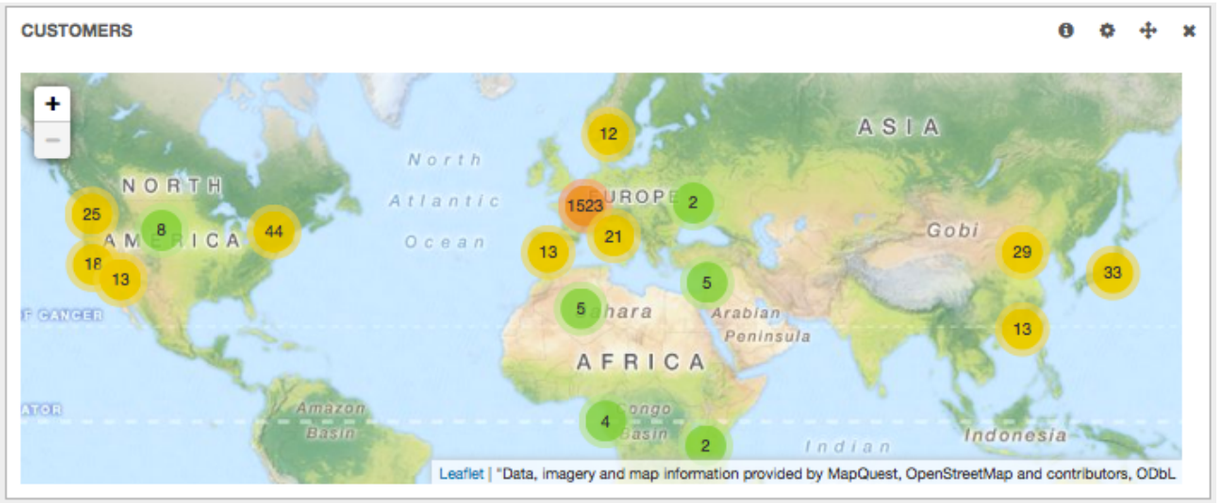

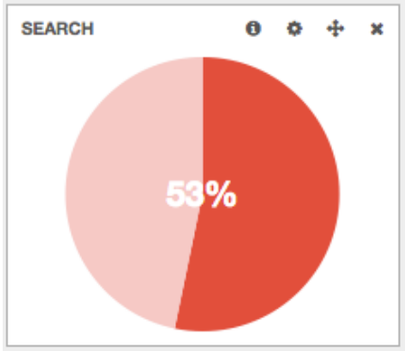

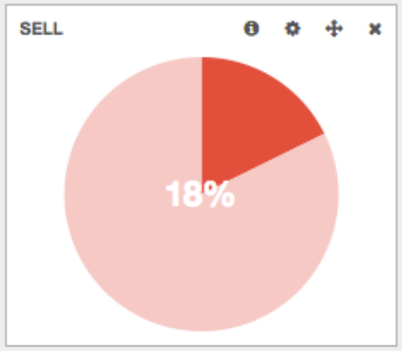

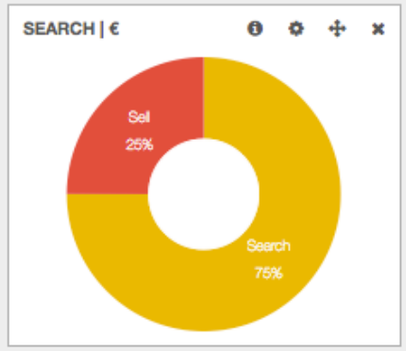

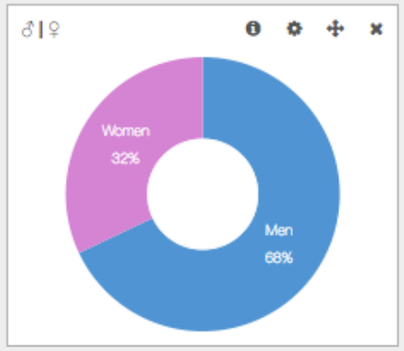

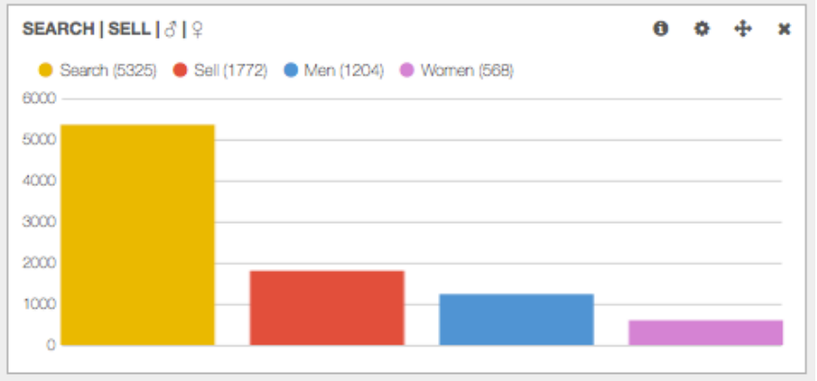

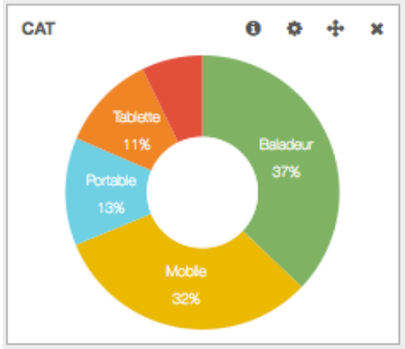

TOP OPTIONS

Term	Count	Action
Disque 32Go	624	
Disque 16Go	582	
Disque 64Go	325	
Disque 2Go	224	
Ecran 13	171	

TOP IP

Term	Count	Action
217.109.123.82	15	
98.223.228.215	9	
95.170.44.138	8	
89.93.253.227	8	
82.225.62.69	8	

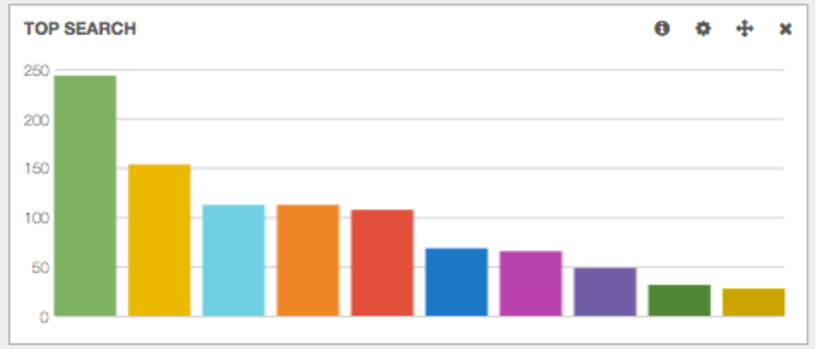

Monitoring Métier / Business

eshop-survey

Journée DEV:OPS

Cluster ELK

Xebia

Elasticsearch – Feedbacks

- The Guardian: social network – real time feedback
- StackOverflow: full-text search with geolocation and « more like »
- Goldman Sacks: 5TB logs/day + analysis stock market
- …

Elasticsearch – NoSQL

SQL	Partitions	DB	Table	Ligne	Colonne
ES	Cluster	Indices	Type	Document	Champ

Elasticsearch – Types de noeuds

- master
- data
- search

Elasticsearch – Shard & Replica

- shards → +indexing, +distribution (one-time setting)
- replicas → +search, +availability

Elasticsearch – Health

- GREEN → all primaries/replicas shards active
- YELLOW → all primaries shards active
- RED → not all primaries shards

Cluster Elasticsearch

http://www.cubrid.org/blog/dev-platform/our-experience-creating-large-scale-log-search-system-using-elasticsearch/

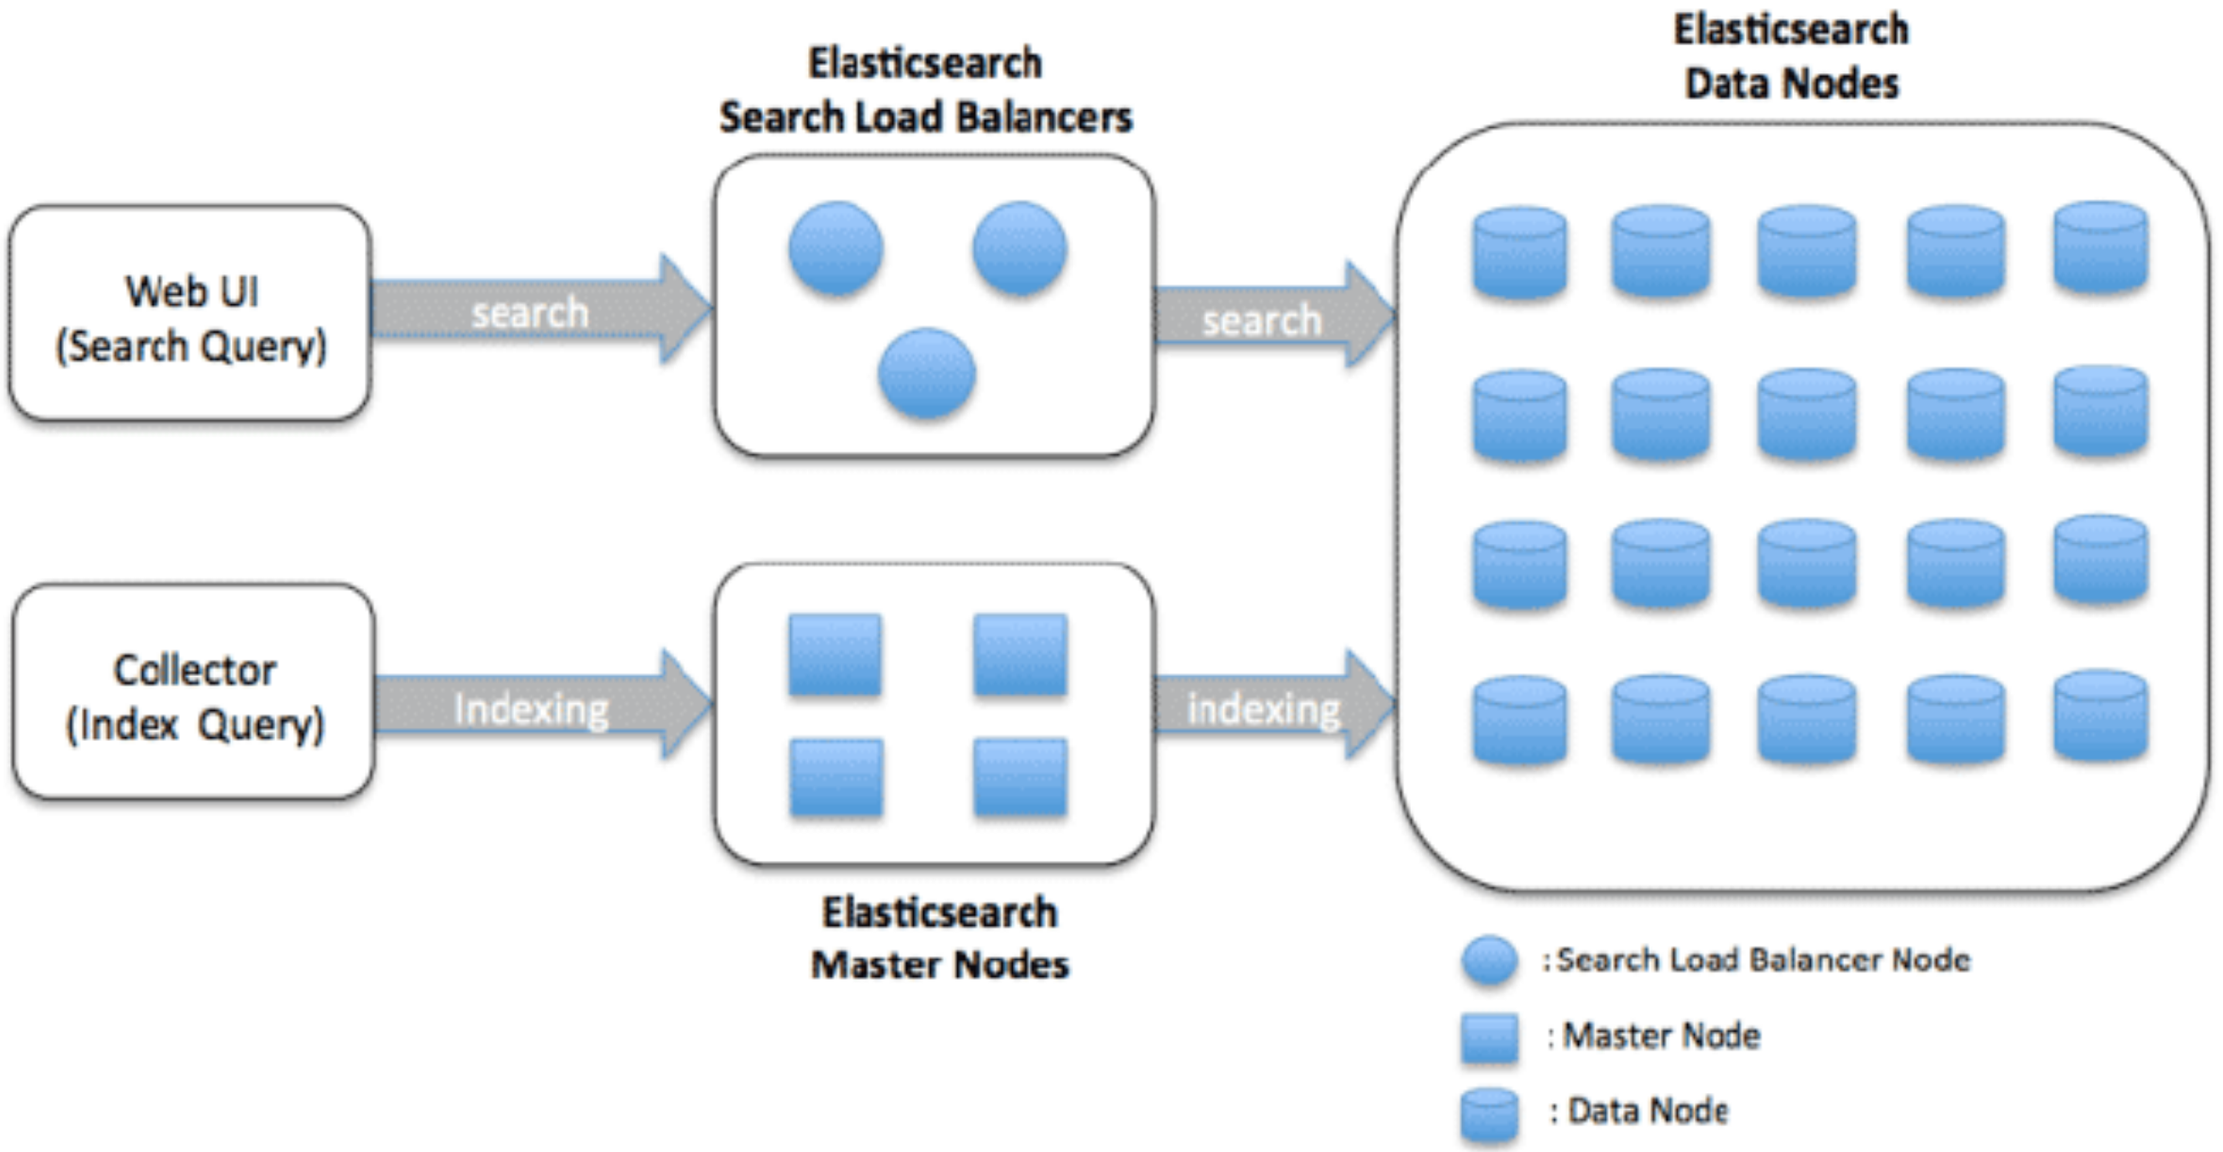

Cluster Elasticsearch

es-cluster

Démo

Journée DEV:OPS

Vagrant (Démos)

Xebia

Démo
Pré-requis

Virtual Box

Vagrant

Git *

Démo @ Home

- https://github.com/vspiewak/elk-devops-day-2014

```
$ tree -L 1 .
.
├── README.md
├── demo-all
├── es-cluster
├── eshop-survey
├── flume-jmx
├── slides
├── syslog
└── system-survey
```

Vagrant – Shortcuts

- cd demo-all

- vagrant up
- vagrant ssh
- sudo jconsole
- vagrant halt*
- vagrant destroy

Vagrant VM

- config.vm.box = "hashicorp/precise64"

- config.vm.network "forwarded_port", guest: 80, host: 10080
- config.vm.network "forwarded_port", guest: 9200, host: 19200

- config.ssh.forward_x11 = true

- vb.customize ["modifyvm", :id, "--ioapic", "on", "--cpuexecutioncap", "40", "--cpus", "2", "--memory", "1024"]

- bootstrap.sh

Questions ?